

# Learning to Rank Places



Shuaib Yunus

## Disclaimer

This thesis was carried out while under the employ of REACH NOW. The dataset used is collected from apps owned by 

## Problem Statement

Solve the ranking problem for geospatial search using Learning to Rank's (LTR) machine learning and deep learning based approaches.



Current Location



haupt



DEPARTURE NOW

STATIONS



Hbf

834 m – Düsseldorf, Deutschland



Hauptbahnhof

7.3 km – Neuss, Deutschland



Hauptfriedhof

9 km – Neuss, Deutschland



Hauptbahnhof

17 km – Solingen, Deutschland



Hbf

19 km – Krefeld, Deutschland

PLACES



Edeka Haupt

7.7 km – Sternstraße 38, 41460 Neuss, Deutschland



Haupt

39 km – Hüttenstraße 47a, 45888 Gelsenkirchen, Deutschland

# Pelias - Existing Search Engine

Based around Elasticsearch:

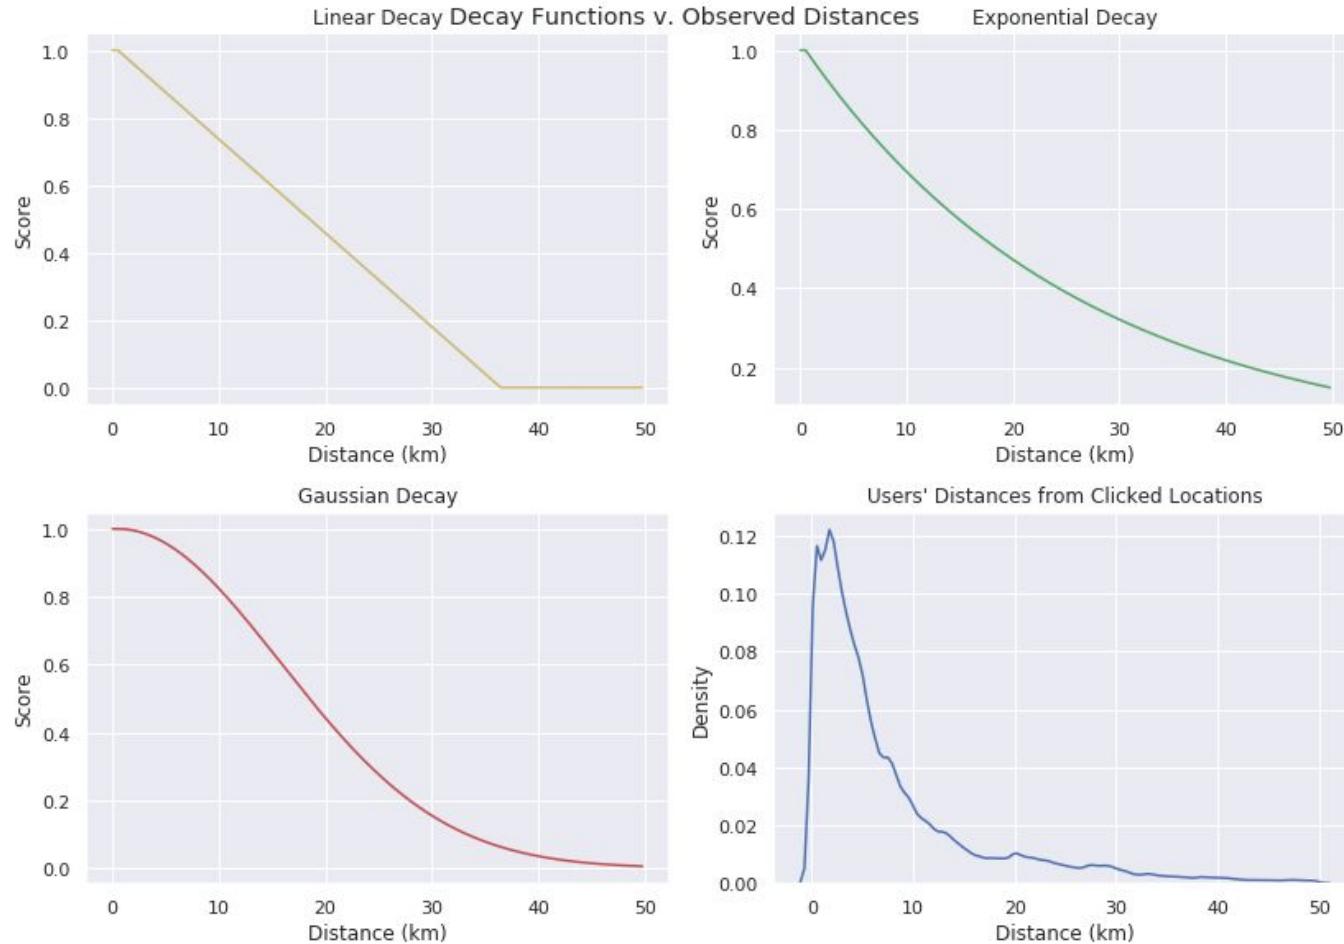
- Performant indexing pipeline for millions of places.
- Sophisticated handwritten analysis chains.
- Intricate Elasticsearch queries to match query texts.

## Limitations of Pelias

Limitations of Pelias are based around Elasticsearch as well.

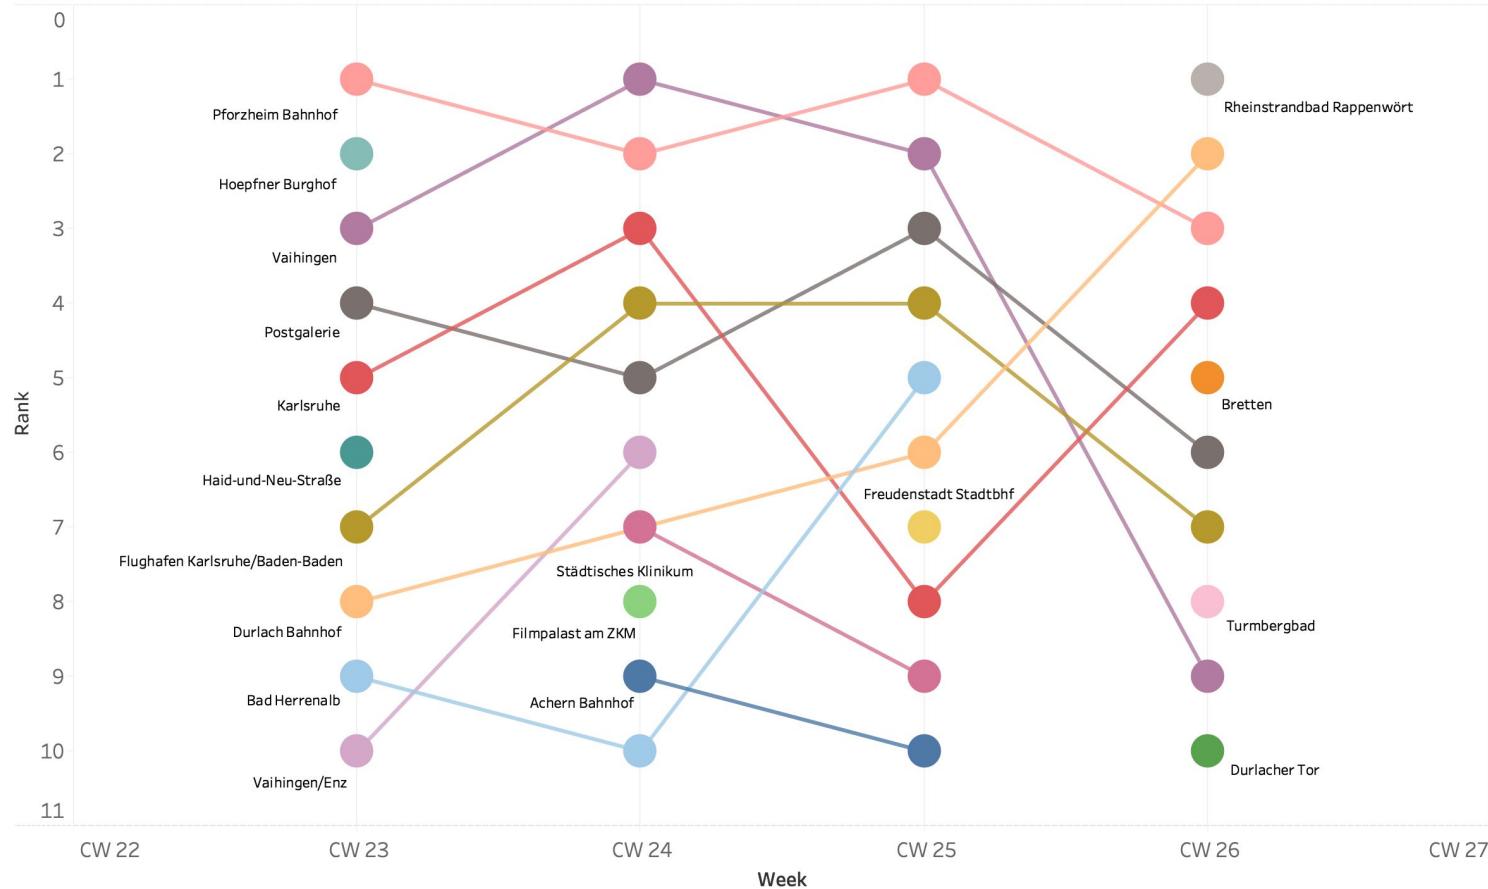
- Learning from Historical Data
- Location Biasing
- Temporal Relevance
- Query Parsing
- Location Sharing

# Location Biasing

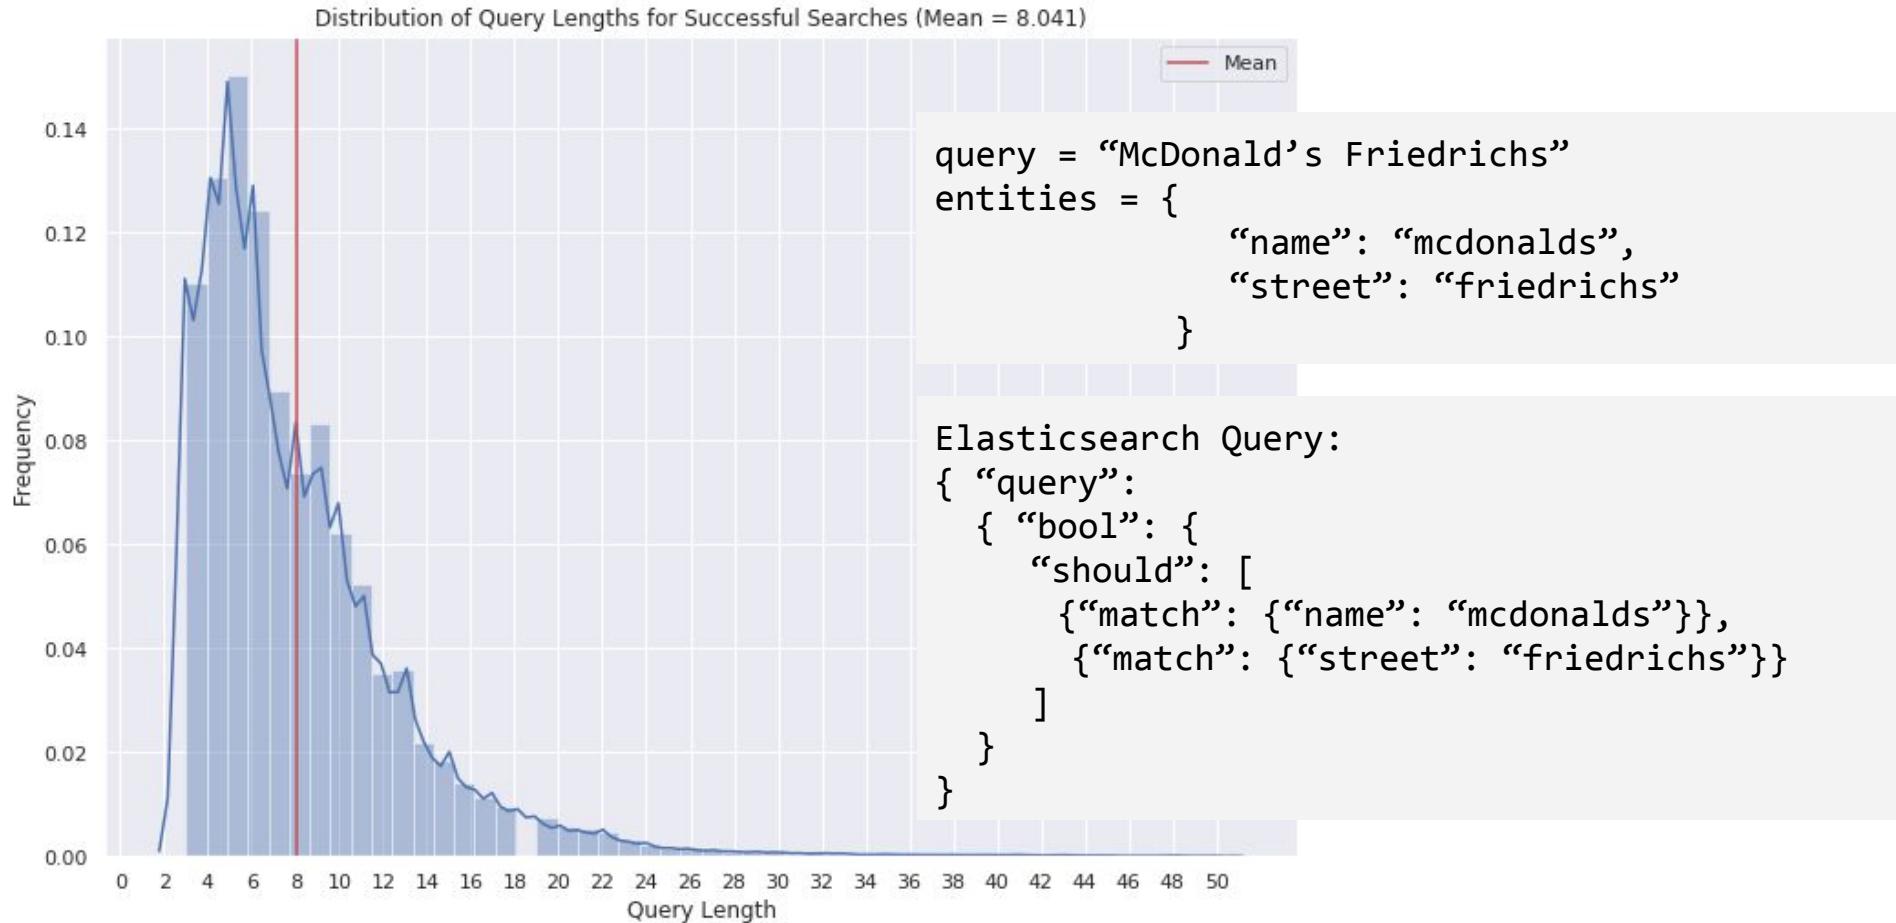


# Temporal Relevance

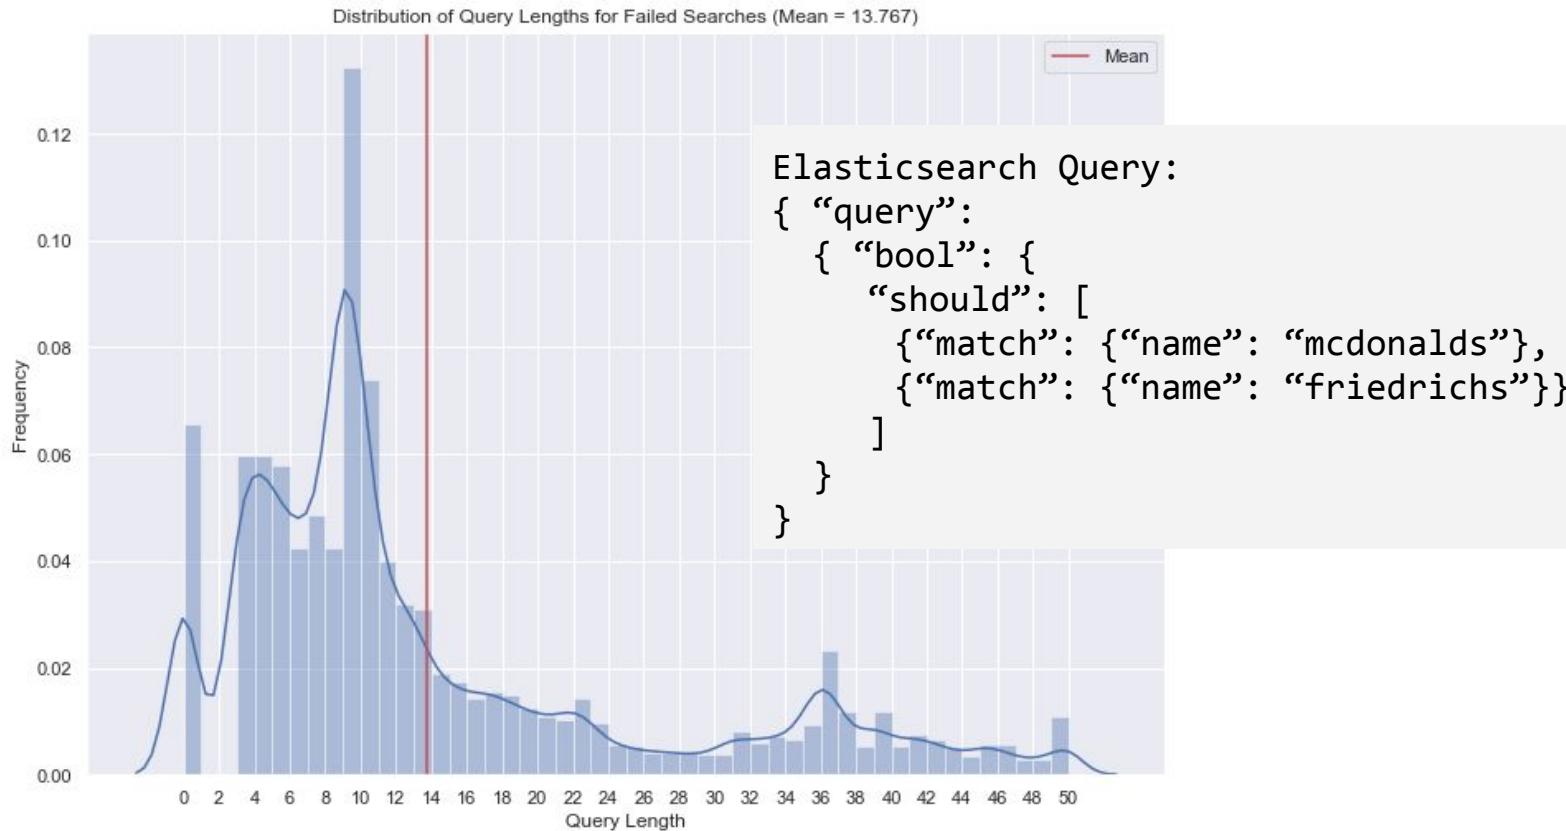
Variations in the ten most selected POIs in Karlsruhe over the weeks of June 2019



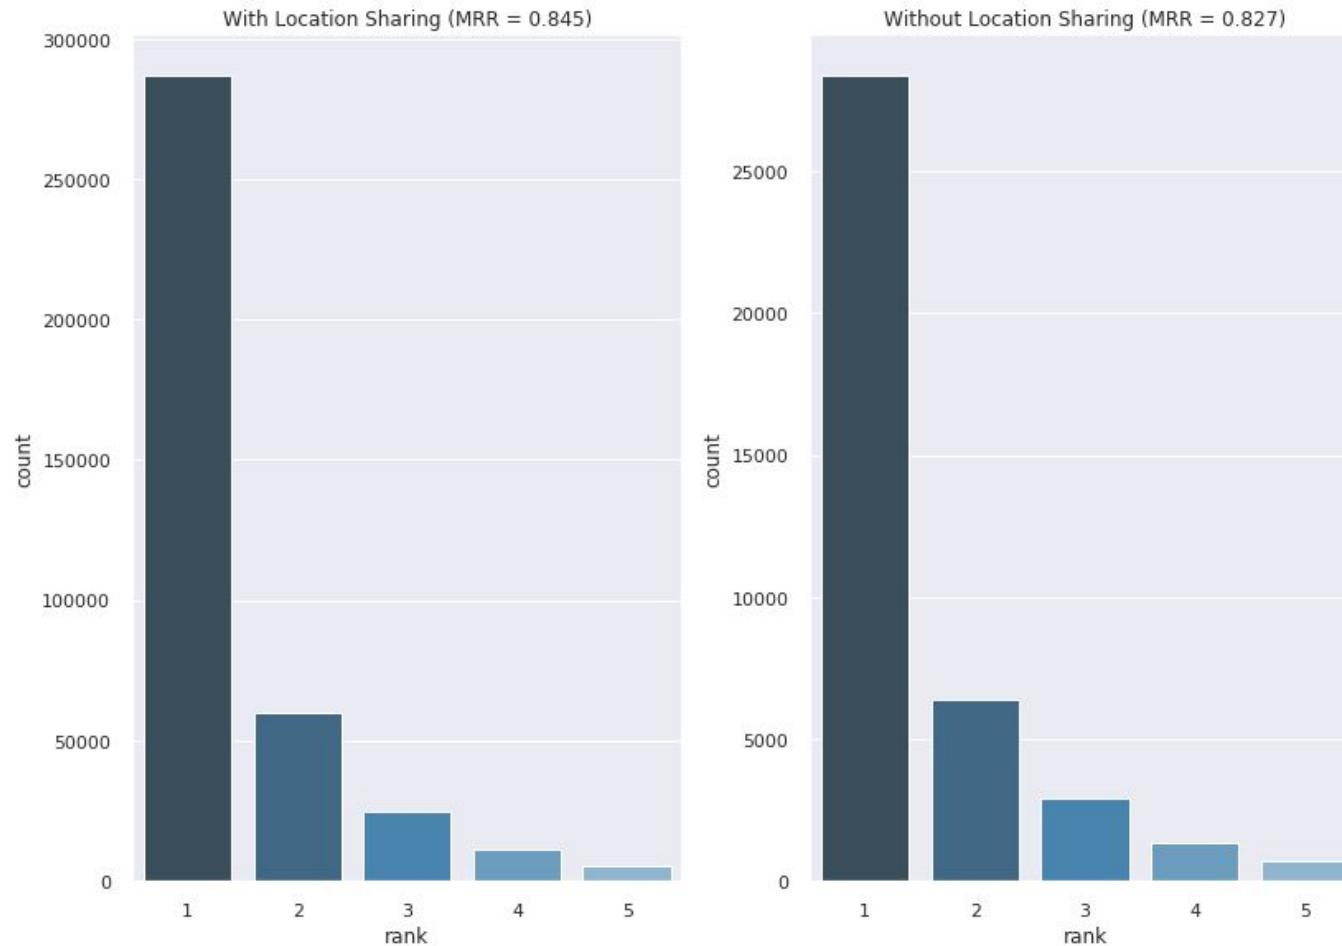
# Query Parsing - Successful Queries are short



# Query Parsing - Failed Searches are lengthy



# Location Sharing



# Learning to Rank (LTR)



## LTR as an Empirical Risk Minimization Problem

The goal is to learn a scoring function  $h : \mathbb{R}^n \rightarrow \mathbb{R}$ , which minimizes:

$$\hat{R}(h) = \frac{1}{n} \sum_{q=1}^n L(\pi(h, X_q), y_q)$$

An LTR algorithm chooses the scoring function  $f$  that minimizes the empirical risk  $\hat{R}(h)$ :

$$f = \arg \min_{h \in \mathcal{H}} \hat{R}(h)$$

The rankings obtained from  $\pi(f, X_q)$  should output the best ordering based on the relevance judgements in the form:

$$y_i^q > y_j^q \Leftrightarrow f(d_i^q) > f(d_j^q)$$

# Dataset

## Metadata

app = KVV.mobil

period = 3 months

search results = 3 mil

no. of selected = 500 k

### Label

### Context Features

### Per-Item Features

Dimension	Description
freq	Number of times the place was selected
Query text	Query text
focus point latitude	Latitude of the user
focus point longitude	Longitude of the user
timestamp	Timestamp of the search
target	Whether the search is for an origin or destination
name	Name of the place
locality	Locality of the place
neighbourhood	Neighbourhood of the presented place
borough	Borough of the place
place latitude	Latitude of the place
place longitude	Longitude of the place
type	Type of the place i.e. poi, address, or station

# Evaluation Metrics

## Mean Reciprocal Rank (MRR)\*

$$MRR = \frac{1}{n} \sum_{i=1}^n \frac{1}{rank_i}$$

## Mean Average Precision at k (MAP@k)

$$\text{MAP}@k = \frac{\sum_{i=1}^U AP_i@k}{n}, \text{ where } AP_i@k = \frac{\sum_{j=1}^{\min\{k, \rho_i\}} rel_{ij} P_i@j}{\sum_{j=1}^{\min\{k, \rho_i\}} rel_{ij}} \quad \& \quad P_i@k = \frac{TP_i}{TP_i + FP_i} = \frac{\sum_{j=1}^{\min\{k, \rho_i\}} rel_{ij}}{k}$$

## Normalized Discounted Cumulative Gain at k (NDCG@k)

$$NDCG_i@k = \frac{DCG_i@k}{IDCG_i@k} \text{ Where, } IDCG_i@k = \sum_{j=1}^{|REL_k|} \frac{2^{rel_j} - 1}{\log_2(j+1)} \quad \& \quad DCG_i@k = \sum_{j=1}^k \frac{2^{rel_j} - 1}{\ln(j+1)}$$

# LTR Approaches

Pointwise à la McRank, Ordinal Regression

$$L(\pi(f, X_q), y_q) = \frac{1}{n} \sum_{i=1}^n (f(d_i^q) - y_i^q)^2$$

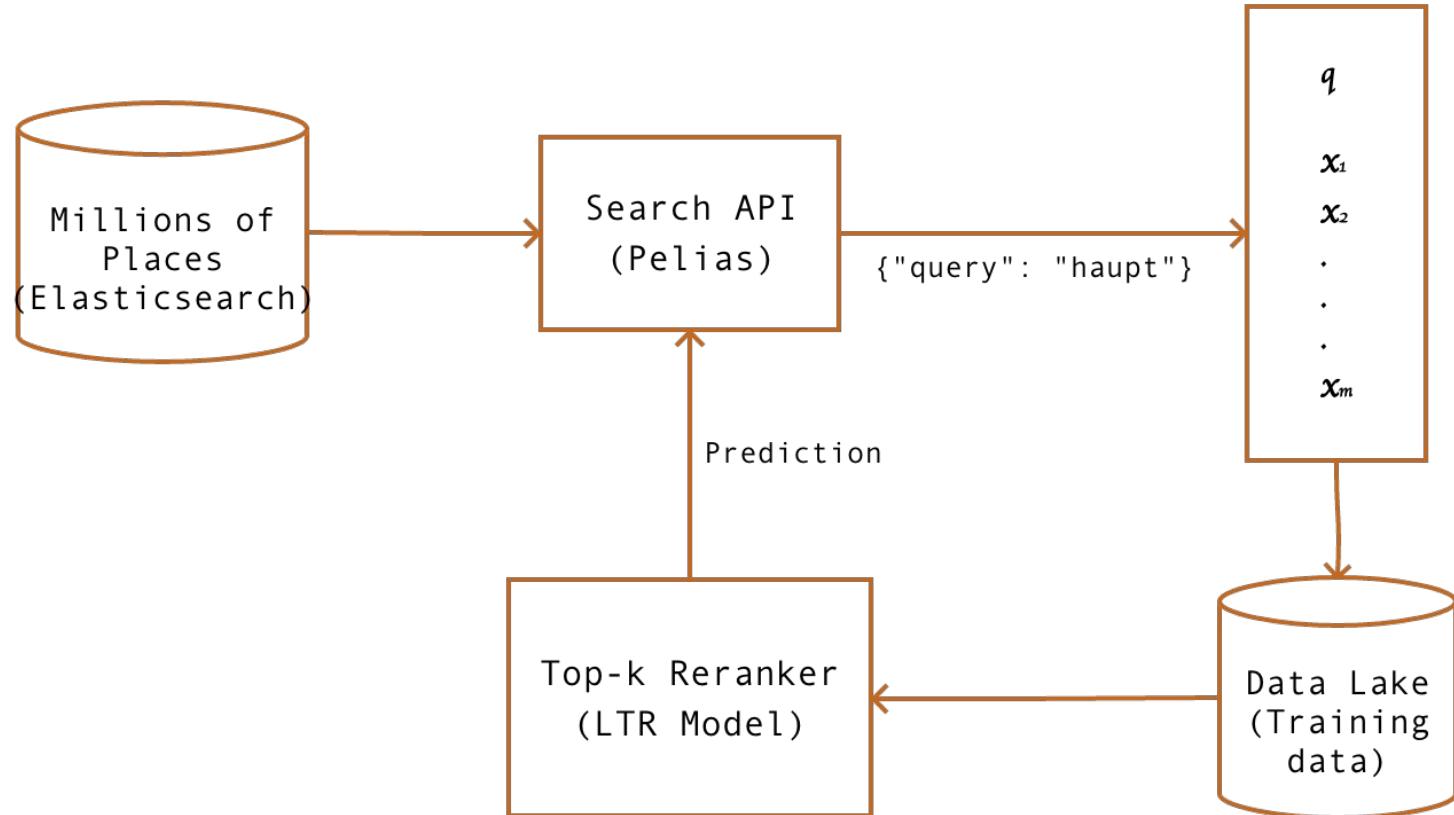
Pairwise à la RankingSVM, RankNet

$$L(\pi(f, X_q), y_q) = \sum_{(i,j): y_i^q < y_j^q} \log \left( 1 + \exp(f(d_i^q)) - \left( f(d_j^q) \right) \right)$$

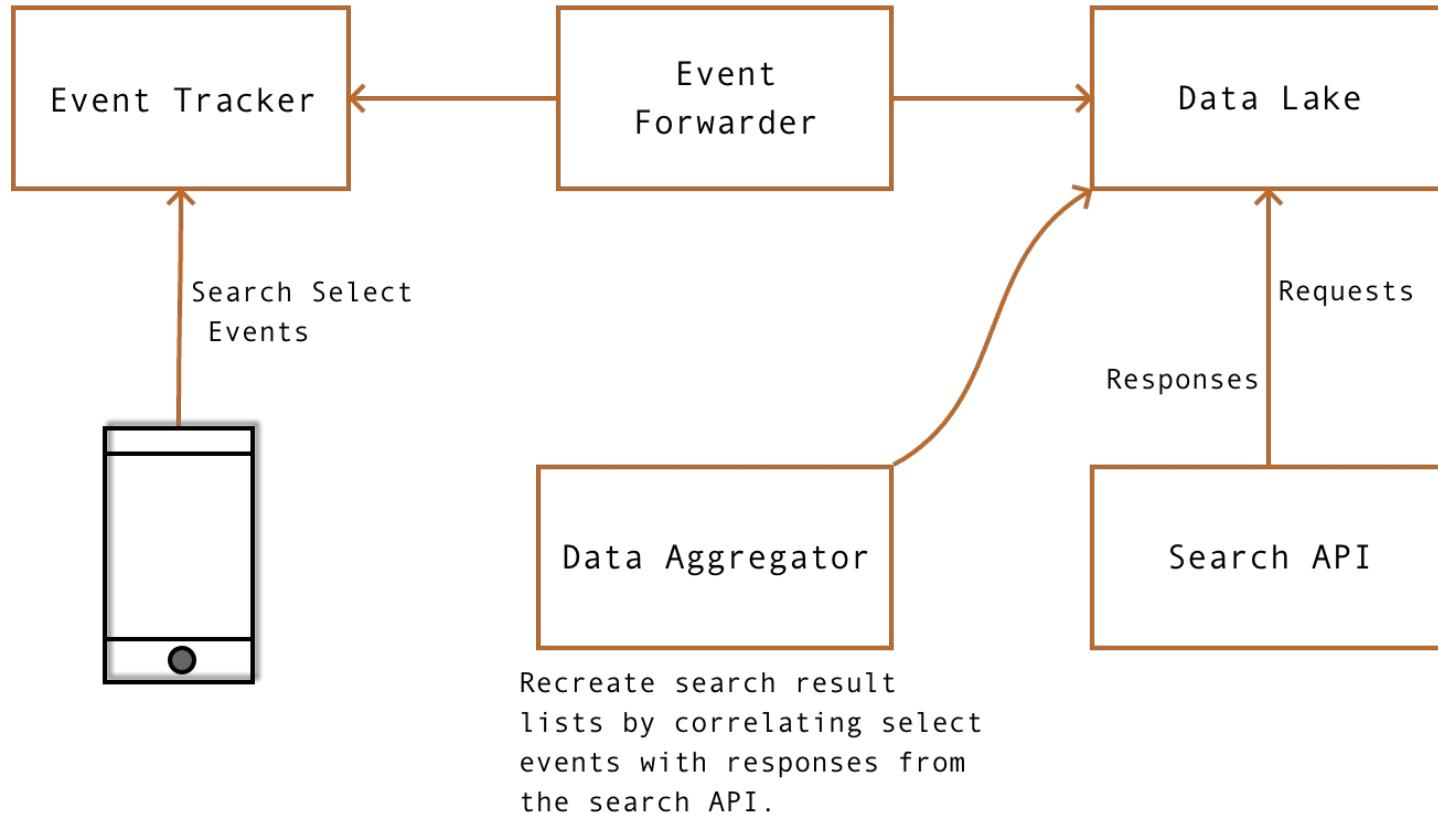
Listwise à la LambdaRank, LambdaMART  $\lambda_{uv} = \frac{-1}{1 + e^{f(x_u) - f(x_v)}}$

LambdaMART and LambdaRank uses pairwise errors, but with weighted gradients based on misranked positions.

# Learning to Rank Framework Used



# Data Engineering



# Feature Engineering

Temporal Features i.e. hour, day, month, year.

$$x_{\sin} = \sin\left(\frac{2 * \pi * x}{\max(x)}\right) \quad x_{\cos} = \cos\left(\frac{2 * \pi * x}{\max(x)}\right)$$

Spatial Features i.e. user's location, place's location.

“lon,lat” => [lon, lat]

Textual Features i.e. query text, place name, city etc.

Hashing vectorizer on character ngrams of range [2,5].

Categorical Features i.e. language, place type (stop, address, poi) etc.

One-hot encoder

# Train-Test-Validation Split

Test = 20%

Validation = 20% of (100% - Test) = 16%

Train = 100% - (Test + Validation) = 64%

Splits treat query groups as a whole, as opposed to the normal split on each observation.

1 query group

 Hbf  
834 m – Düsseldorf, Deutschland

 Hauptbahnhof  
7.3 km – Neuss, Deutschland

 Hauptfriedhof  
9 km – Neuss, Deutschland

 Hauptbahnhof  
17 km – Solingen, Deutschland

 Hbf  
19 km – Krefeld, Deutschland

## PLACES

 Edeka Haupt  
7.7 km – Sternstraße 38, 41460 Neuss, Deutschland

 Haupt  
39 km – Hüttenstraße 47a, 45888 Gelsenkirchen, Deutschland

# Tree-Based Ranker

ML Algorithm: Gradient Boosted Trees

LTR Algorithms: LambdaRank, LambdaMART

Libraries: XGBoost, LightGBM

Hyperparameters:

- Objective function i.e. MAP, NDCG, Pairwise
- Scaling of positive weights i.e. selected results
- Number of rounds i.e. number of trees.
- Minimum child weight
- Maximum tree depth

# Neural Ranker

ML Algorithm: Neural Network

LTR Algorithms: LambdaRank, LambdaMART

Libraries: TensorFlow Ranking

Optimizations: Multi-Item Scoring (Groupwise Scoring Functions),  
Ranking Metric Optimization (LambdaLoss)

Hyperparameters:

- NN Architecture i.e. activation fn, hidden layers, etc.
- LambdaLoss Metric i.e. MRR, NDCG, Mean Squared, etc.
- Group size for Groupwise Scoring Functions

# Results

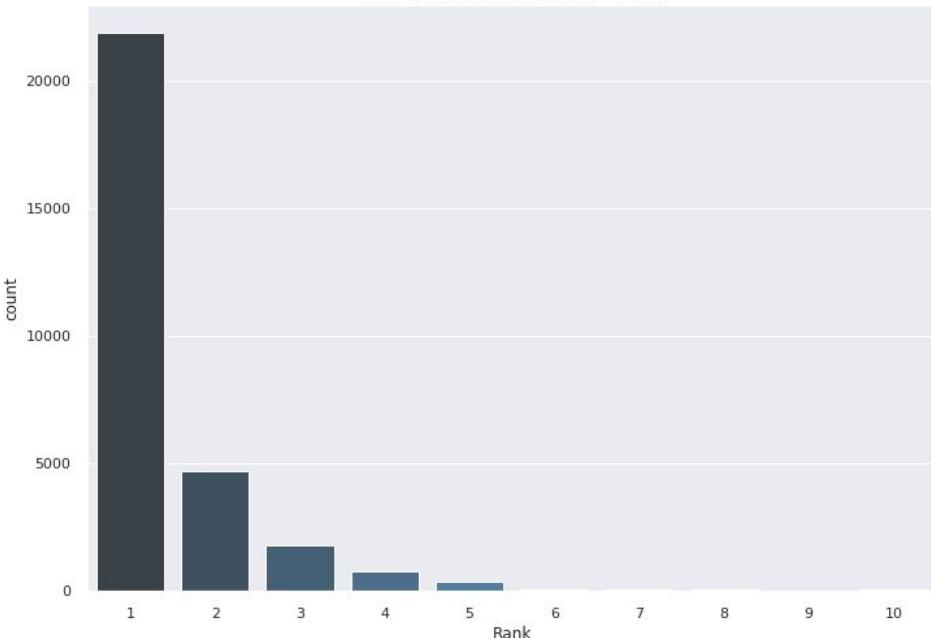


# Model Comparisons

Model	MRR	
<i>Pelias</i>	<b>0.8493</b>	-
SVM <sup>rank</sup>	0.7415	<b>-13%</b>
XGBoost	0.8754	+3.1%
<b>LightGBM</b>	<b>0.8922</b>	<b>+5.1%</b>
TF-Ranking (Linear)	0.8675	+2.2%
TF-Ranking (Deep)	0.8559	+0.78%

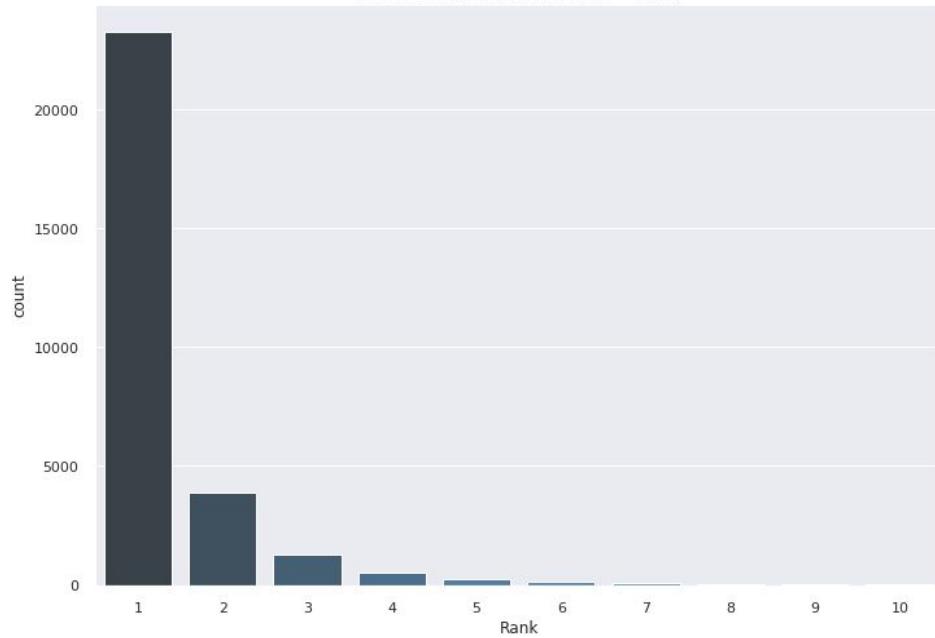
# Rank Distributions

Pelias Rank Distribution (MRR = 0.849)



Pelias

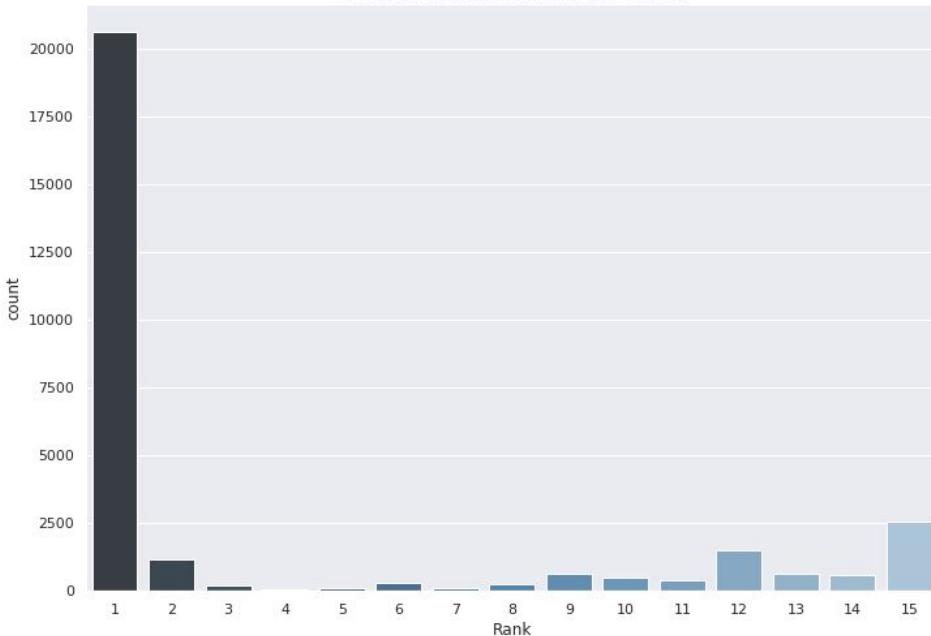
XGBoost Rank Distribution (MRR = 0.875)



XGBoost

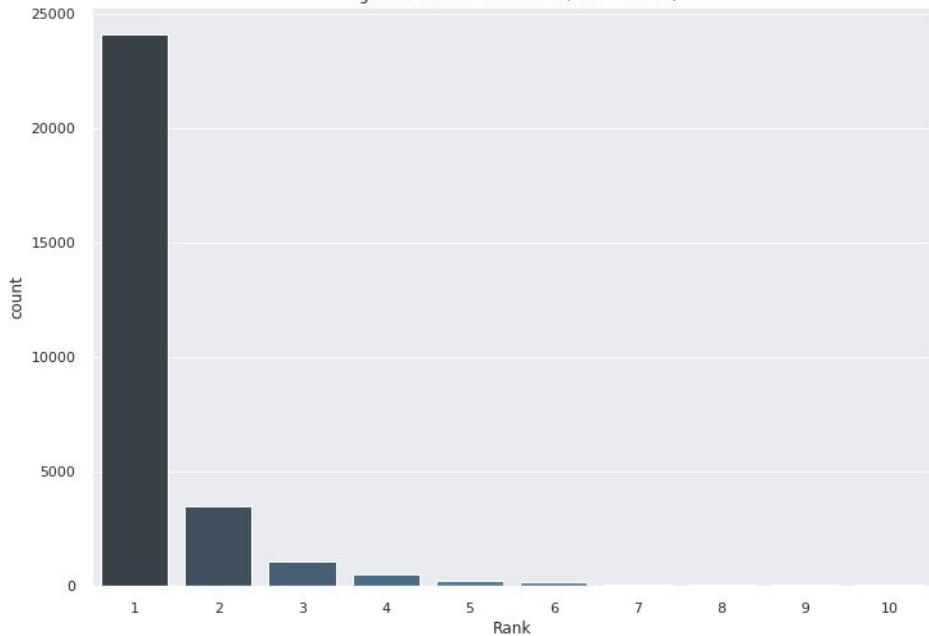
# Rank Distributions (cont.)

SVRMRank Rank Distribution (MRR = 0.742)



SVM<sup>rank</sup>

LightGBM Rank Distribution (MRR = 0.892)



LightGBM

# Tree-Based Ranker - Hyperparameter Tuning

Model	MRR
XGBoost (Default)	0.8531
<b>XGBoost (HP Tuned)</b>	<b>0.8754</b>
LightGBM (Default)	0.8826
<b>LightGBM (HP Tuned)</b>	<b>0.8922</b>

# Tree-Based Feature Importances

Top 20 non-textual features

Weight	Feature
0.1574	x0_stop
0.0036	x0_address
0.0010	x0_poi
0.0005	place_lon
0.0005	place_lat
0.0003	focus_lat
0.0002	focus_lon
0.0000	month_sin
0.0000	hr_cos
0.0000	hr_sin
0.0000	day_sin
0.0000	is_weekend
0.0000	day_cos
0	month_cos

Top 20 features

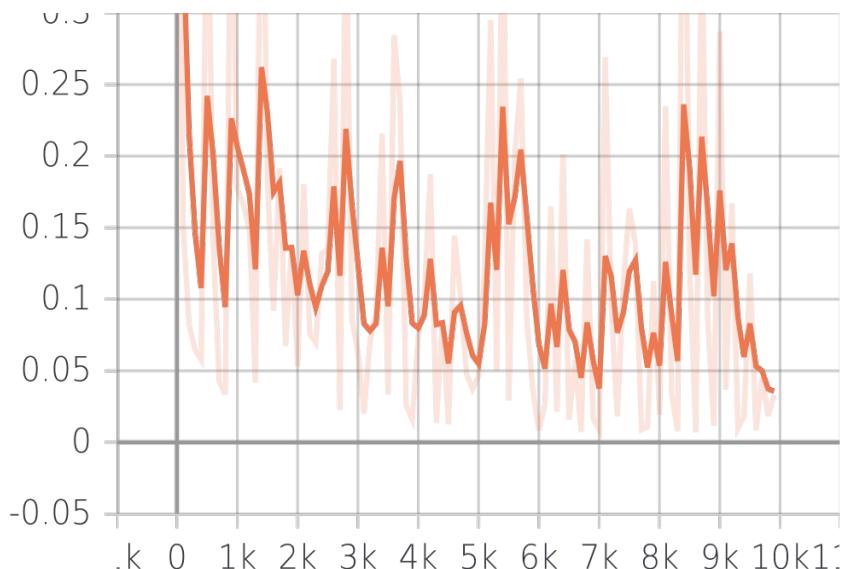
Weight	Feature
--------	---------

# Neural Ranker Architectures

#	Model	MRR
1	Linear (sigmoid) (0 hidden layers)	0.8675
2	FC-64-32-16: 80% dropout	0.8559
3	FC-32 (shallow network, 80% dropout)	0.8523
4	FC-512-256-32 (groupwise with size=5)	0.8247
5	FC-512-256-32	0.8070
6	FC-512-256-32 (tanh)	0.8039
7	FC-64-32-16	0.8038
8	FC-512-256-32 (approximate MRR loss)	0.7790
9	Non-Linear (ReLU) (0 hidden layers)	0.7732

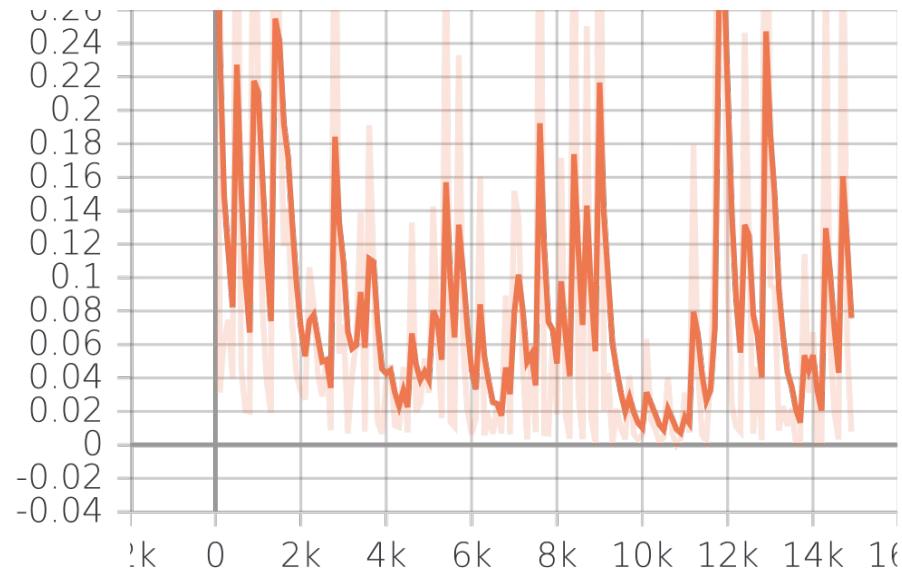
# Neural Ranker - Training Loss

**Linear (Sigmoid)**



**MRR = 0.8675**

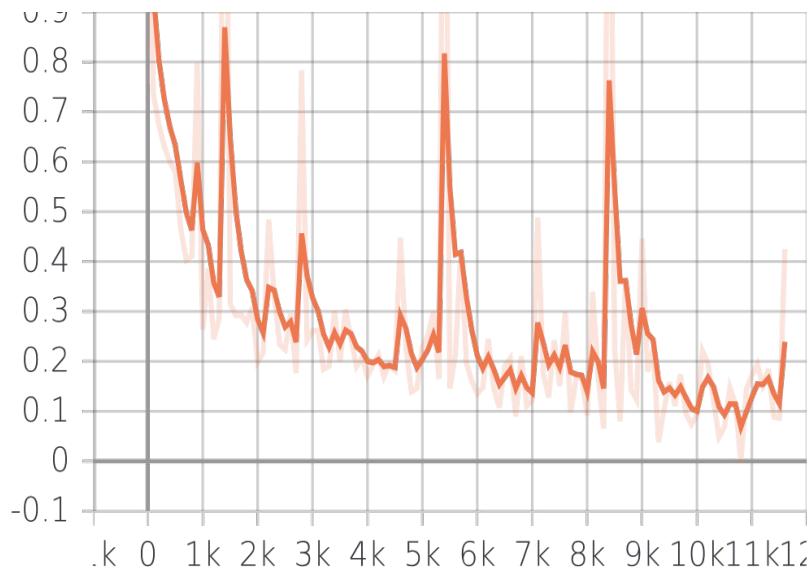
**FC-512-256-32**



**MRR = 0.8070**

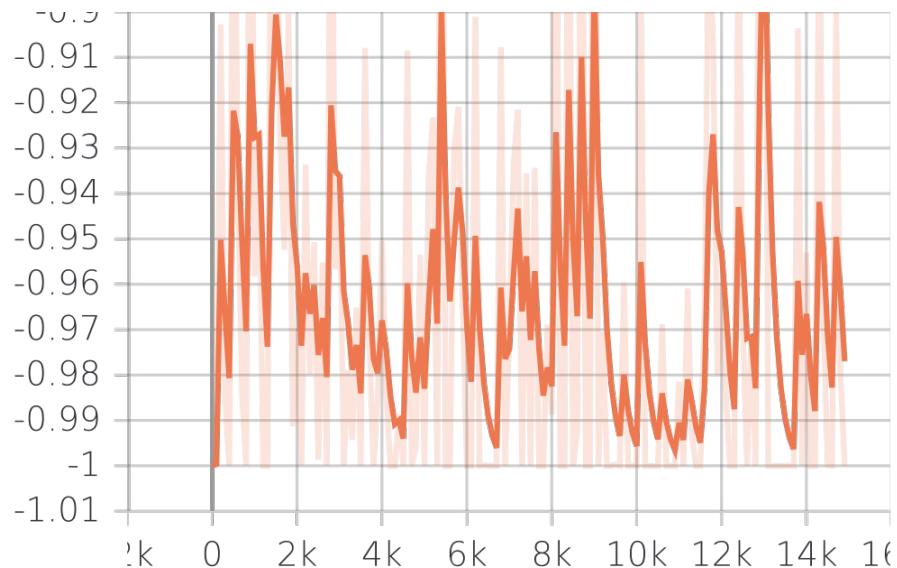
# Neural Ranker - Training Loss

**FC-64-32-16 (80% Dropout)**



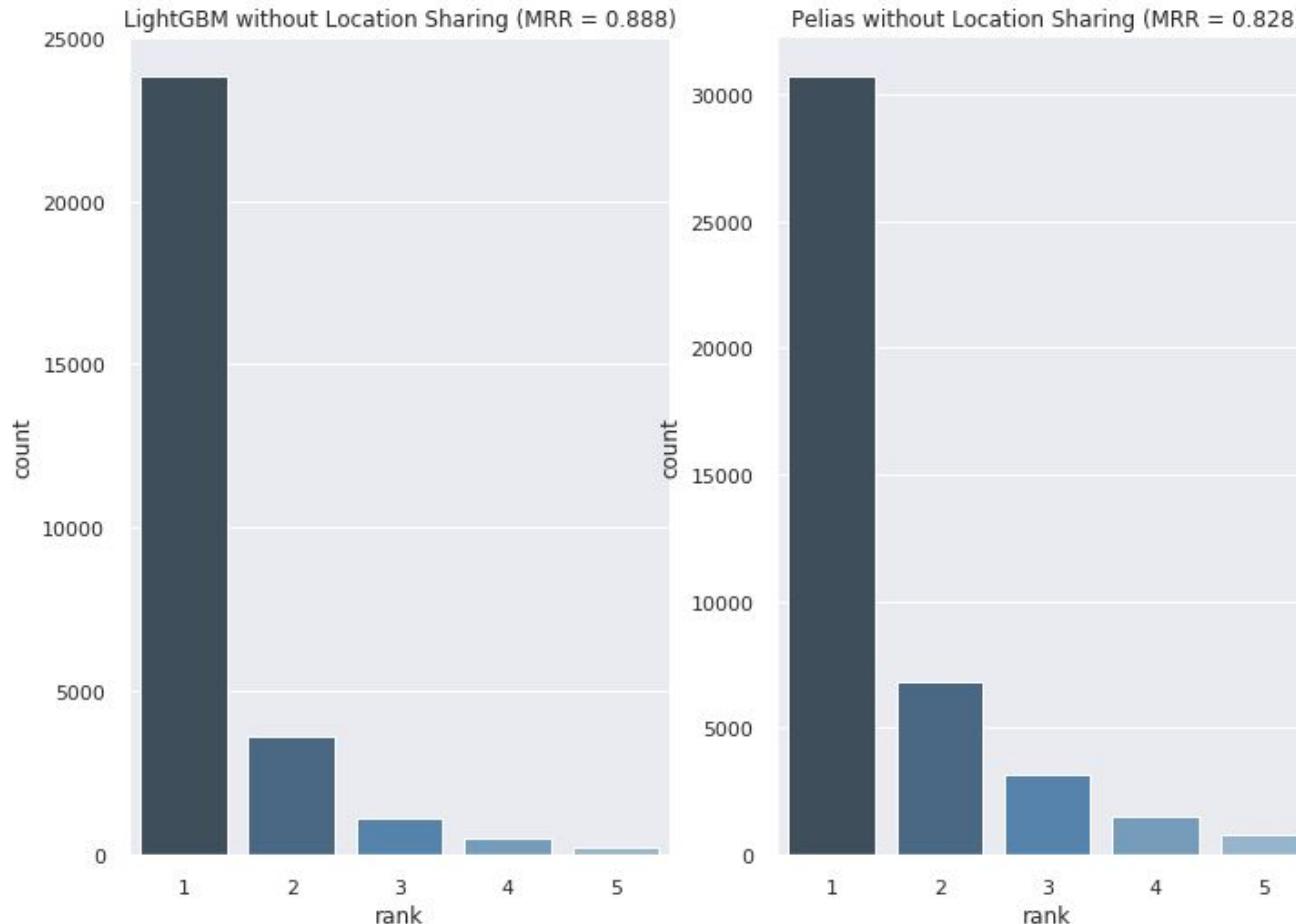
**MRR = 0.8559**

**FC-512-256-32 (MRR Loss)**



**MRR = 0.7790**

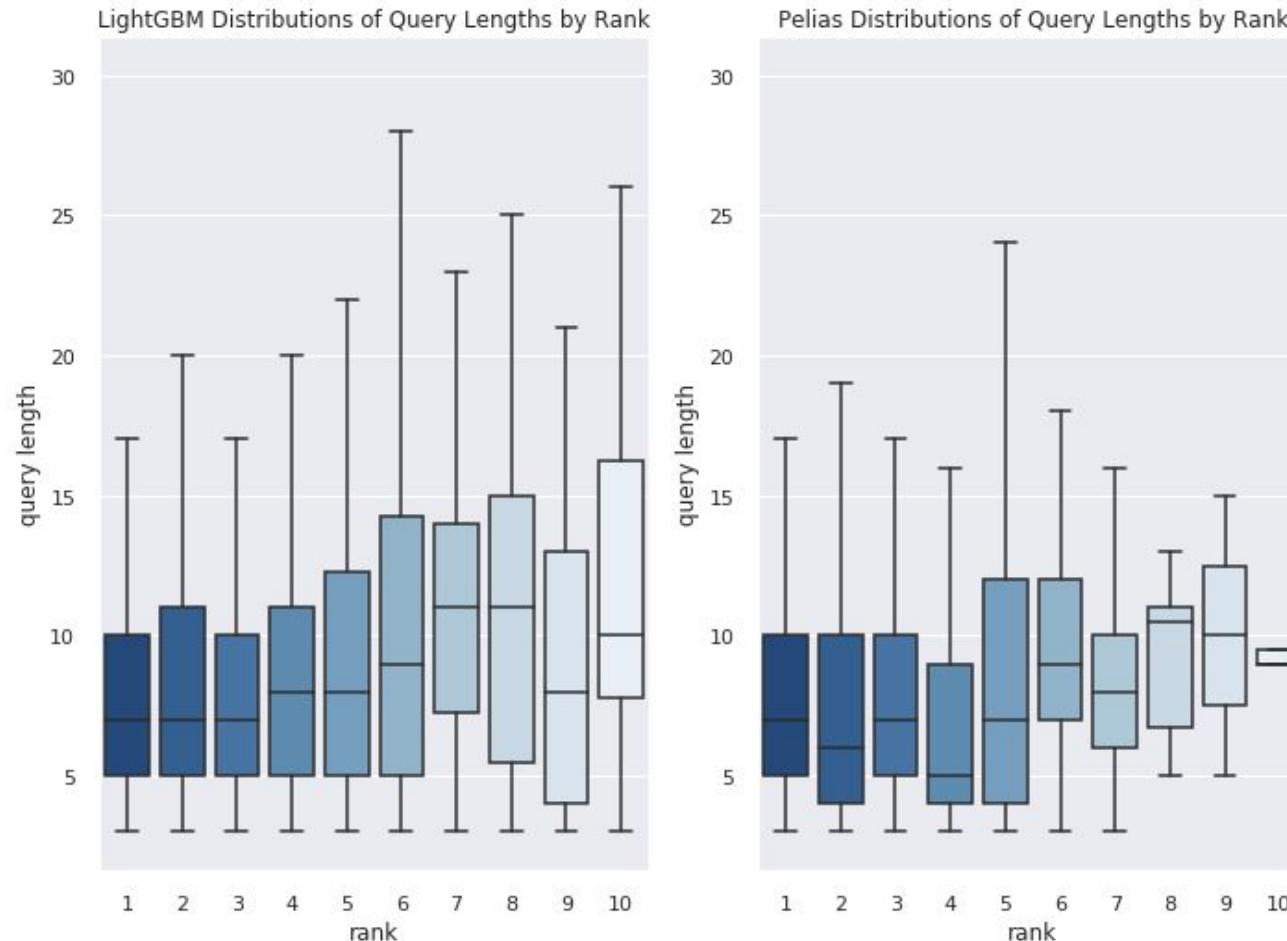
# Model vs Pelias - Location Sharing



# Model vs Pelias - Query Parsing

Query Character Length Range	LightGBM MRR	Pelias MRR
[3, 10]	0.9003	0.8481
[11, 20]	0.8721	0.8531
[21, 30]	0.8044	0.8612
[31, 40]	0.5647	0.8167

# Model vs Pelias - Query Parsing (cont.)



# Tree-Based Rankers vs Neural Rankers

	Tree-Based	Neural
Training Time		
Hyperparameter Tuning		
Ease of implementation		
Cool factor		

# Conclusion

- LTR models show improvement as top-k rerankers on top of Pelias.
- Tree-based rankers produced higher MRR scores than neural rankers.
- Both tree-based and neural rankers reveal potential for significant improvements on modest tuning.

## Future Research

- Interaction between Pelias and the LTR rankers to optimize for different respective retrieval phases.
- Explore applicability of unbiased learning to rank methods.
- Incorporating unsuccessful queries during training.
- Rankers for worldwide search using city-based datasets.



# Thank you!